

CS 435 Parallel and Distributed Processing

Project Report

High Performance Computing

M. Saad Farrukh

Reg. No: 373188

Section: BEE-13 D

Instructor: Dr. Attique Dawood

11:48 PM PKT, Saturday, May 24, 2025

Contents

1	Part I: Getting to Know the HPC Cluster	3
1.1	Connecting to the Server	3
1.2	Checking Reachable/Active Nodes	3
1.2.1	Manual Method	3
1.2.2	Scripted Method	3
1.3	CPU Specifications	5
1.3.1	Manual Method	5
1.3.2	Scripted Method	5
1.4	Checking Busy Nodes	6
1.4.1	Manual Method	6
1.4.2	Scripted Method	6
1.5	Ranking Nodes	8
2	Part II: Laplace Solver Implementation	9
2.1	Implementation Details	9
2.2	Performance Testing with 4 Nodes	10
2.2.1	High Node = 4	10
2.2.2	Contour Plots	12
2.3	Performance Testing with 3 Nodes	13
2.3.1	Node = 3	13
2.3.2	Contour Plots	15
2.4	Comparison of 4 and 3 Nodes	16
2.5	Performance Bottlenecks and Comparative Analysis	16
2.5.1	Why MPI4 Outperformed MPI3	17
2.5.2	Why MPI5 Performed Poorly	17
2.6	Additional : Data Transfer	17
A	Full MPILaplace Code	18
A.1	MPILaplace.cpp (Abbreviated)	18
A.2	Full MPILaplace.cpp	18

1 Part I: Getting to Know the HPC Cluster

This section details the analysis of the SINES HPC cluster, conducted on Saturday, May 17, 2025, at approximately 9:00 PM.

The analysis aimed to map compute nodes, determine CPU specifications, identify busy nodes, and rank nodes by performance. The cluster includes testing nodes `compute-0-1` to `compute-0-50`, with 16 nodes confirmed available.

1.1 Connecting to the Server

Access was established using:

1. VPN credentials from SEECs Support were configured.
2. Connection command:

```
1 ssh -X user10@10.19.10.50 -oHostKeyAlgorithms=+ssh-rsa
```

Technical Insight: The `-X` flag enables X11 forwarding for graphical output, while `-oHostKeyAlgorithms=+ssh-rsa` supports legacy SSH keys, critical for cluster access.

1.2 Checking Reachable/Active Nodes

Node availability was assessed to identify operational units.

1.2.1 Manual Method

Manually running `ssh compute-0-X` for each node was infeasible due to the scale.

Although tried that approach but it was taking much time, with large number of screenshots.

1.2.2 Scripted Method

A bash script automated the process:

```
1 count=0
2 for i in {1..50}; do
3     if ssh -o ConnectTimeout=1 compute-0-$i "exit" 2>/dev/null;
4         then
5             echo "compute-0-$i is up"
6             ((count++))
7         fi
8 done
9 echo "Total available nodes: $count"
```

Output: 16 nodes were available, found to be up..

Table 1: Available Compute Nodes

Node	Status
compute-0-5	Up
compute-0-6	Up
compute-0-8	Up
compute-0-9	Up
compute-0-10	Up
compute-0-11	Up
compute-0-12	Up
compute-0-13	Up
compute-0-16	Up
compute-0-17	Up
compute-0-18	Up
compute-0-19	Up
compute-0-23	Up
compute-0-24	Up
compute-0-26	Up
compute-0-27	Up

```
[user10@afrit ~]$ nano check_nodes.sh
[user10@afrit ~]$ chmod +x check_nodes.sh
[user10@afrit ~]$ ./check_nodes.sh
compute-0-5 is up
compute-0-6 is up
compute-0-8 is up
compute-0-9 is up
compute-0-10 is up
compute-0-11 is up
compute-0-12 is up
compute-0-13 is up
compute-0-16 is up
compute-0-17 is up
compute-0-18 is up
compute-0-19 is up
compute-0-23 is up
compute-0-24 is up
compute-0-26 is up
compute-0-27 is up
user10@compute-0-31's password:
Total available nodes: 16
[user10@afrit ~]$
```

Figure 1: Reachable Nodes - Total 16 reported

Technical Insight: The `ConnectTimeout=1` ensures rapid failure detection, optimizing script runtime.

1.3 CPU Specifications

CPU details were collected for performance evaluation.

1.3.1 Manual Method

Running `lscpu` manually on each node was time-intensive.

1.3.2 Scripted Method

A bash script gathered data:

```

1 #!/bin/bash
2 printf "%-15s %-10s %-15s %-10s %-10s %-12s %-10s\n" "Node" "CPU(
   s)" "Architecture" "CPU MHz" "BogoMIPS" "NUMA node(s)" "Socket(
   s)"
3 for i in {1..50}; do
4     ssh -o ConnectTimeout=1 compute-0-$i 'lscpu 2>/dev/null |
       grep -E "^CPU\(s\):|Architecture|CPU MHz|BogoMIPS|NUMA node
       \(s\)|Socket\(s\)" | awk -F: "{print \$2}" | xargs || echo
       "N/A"' 2>/dev/null
5 done

```

Table 2: CPU Specifications

Node	CPU(s)	Architecture	CPU MHz	BogoMIPS	NUMA Node(s)
compute-0-5	16	x86_64	1600.000	4532.70	2
compute-0-6	16	x86_64	2667.000	4532.69	2
compute-0-8	24	x86_64	1600.000	5333.22	2
compute-0-9	16	x86_64	2667.000	4532.71	2
compute-0-10	16	x86_64	1600.000	4532.70	2
compute-0-11	16	x86_64	1600.000	4532.70	2
compute-0-12	16	x86_64	1600.000	4532.70	2
compute-0-13	16	x86_64	1600.000	4532.71	2
compute-0-16	16	x86_64	1600.000	4532.70	2
compute-0-17	16	x86_64	1600.000	4532.72	2
compute-0-18	16	x86_64	1600.000	4532.70	2
compute-0-23	16	x86_64	2667.000	4532.69	2
compute-0-24	16	x86_64	1600.000	4532.69	2
compute-0-26	16	x86_64	1600.000	4532.69	2
compute-0-27	16	x86_64	1600.000	4532.70	2

Technical Insight: `lscpu` confirms uniform hardware (16 cores, except `compute-0-6` with 24.) Few compute nodes seems to be overclocked.

```
[user10@afrit ~]$ nano table_nodes.sh
[user10@afrit ~]$ chmod +x table_nodes.sh
[user10@afrit ~]$ ./table_nodes.sh
```

Node	CPU(s)	Architecture	CPU MHz	BogoMIPS	NUMA Node(s)	Socket(s)
compute-0-5	16	x86_64	1600.000	4532.70	2	2
compute-0-6	16	x86_64	2267.000	4532.69	2	2
compute-0-8	24	x86_64	1600.000	5333.22	2	2
compute-0-9	16	x86_64	2267.000	4532.71	2	2
compute-0-10	16	x86_64	1600.000	4532.70	2	2
compute-0-11	16	x86_64	1600.000	4532.70	2	2
compute-0-12	16	x86_64	1600.000	4532.70	2	2
compute-0-13	16	x86_64	1600.000	4532.71	2	2
compute-0-16	16	x86_64	1600.000	4532.70	2	2
compute-0-17	16	x86_64	2267.000	4532.72	2	2
compute-0-18	16	x86_64	1600.000	4532.70	2	2
compute-0-23	16	x86_64	2267.000	4532.69	2	2
compute-0-24	16	x86_64	1600.000	4532.69	2	2
compute-0-26	16	x86_64	1600.000	4532.69	2	2
compute-0-27	16	x86_64	1600.000	4532.70	2	2

Figure 2: Table Nodes - CPU specs for available compute nodes

1.4 Checking Busy Nodes

Busy nodes were identified using resource metrics.

1.4.1 Manual Method

Using `top` manually was inefficient.

1.4.2 Scripted Method

A script collected data:

```
1 printf "%-15s %-10s %-20s %-10s\n" "Node" "Mem%" "Load Avg" "
   Buffers"
2 for i in {1..50}; do
3     ssh -o ConnectTimeout=1 compute-0-$i 'free -m 2>/dev/null |
       awk "/Mem:/ {printf \"%.1f\\\", (\\$3/\\$2)*100}" && uptime 2>/
       dev/null | awk -F "load average:" "{print \\$2}" | awk "{
       printf \"%.2f\\\", \\$1}"' 2>/dev/null
4 done
```

Technical Insight: `compute-0-23` shows high memory usage (81.4%) and load (18.00) and `compute-0-11` shows low memory usage (4.3%) and load (0.00), indicating resource contention.

Variability suggests dynamic workloads.

Table 3: Node Busyness

Node	Mem%	Load Avg	Buffers
compute-0-5	14.0	0.00	162MB
compute-0-6	21.5	15.13	139MB
compute-0-8	25.3	0.00	144MB
compute-0-9	52.0	12.00	218MB
compute-0-10	39.6	0.05	205MB
compute-0-11	4.3	0.00	203MB
compute-0-12	45.7	7.14	213MB
compute-0-13	28.2	12.22	198MB
compute-0-16	17.1	0.00	191MB
compute-0-17	35.6	18.04	220MB
compute-0-28	10.0	0.00	197MB
compute-0-23	81.4	18.00	211MB
compute-0-24	6.0	0.04	206MB
compute-0-26	3.5	0.00	200MB
compute-0-27	18.2	0.00	237MB

```

[user10@afrit ~]$ nano performance_nodes.sh
[user10@afrit ~]$
[user10@afrit ~]$ chmod +x performance_nodes.sh
[user10@afrit ~]$ ./performance_nodes.sh
Node      Mem%      Load Avg      Buffers
compute-0-5  14.0%     0.00           162MB
compute-0-6  21.5%     15.13          139MB
compute-0-8  25.3%     0.00           144MB
compute-0-9  52.0%     12.00          218MB
compute-0-10 39.6%     0.05           205MB
compute-0-11 4.3%      0.00           203MB
compute-0-12 45.7%     7.14           213MB
compute-0-13 28.2%     12.22          198MB
compute-0-16 17.1%     0.00           191MB
compute-0-17 35.6%     18.04          220MB
compute-0-18 10.0%     0.00           197MB
compute-0-23 81.4%     18.00          211MB
compute-0-24 6.0%      0.00           206MB
compute-0-26 3.5%      0.00           200MB
compute-0-27 18.2%     0.00           237MB

```

Figure 3: Performance Nodes - Memory%, Load Avg, Buffers for Available Compute Nodes

1.5 Ranking Nodes

Nodes were ranked using:

$$\text{Score} = \left(\frac{1}{\text{Mem}\% + 0.01} \right) \times \left(\frac{1}{\text{Load Avg} + 0.01} \right) \times \left(\frac{\text{Buffers (MB)}}{100} \right)$$

(Used 0.01 for zero load averages.)

Table 4: Node Performance Ranking

Node	Score	Rank
compute-0-13	1.312	1
compute-0-26	1.247	2
compute-0-11	1.207	3
compute-0-24	1.150	4
compute-0-16	1.149	5
compute-0-5	1.075	6
compute-0-6	1.058	7
compute-0-8	1.046	8
compute-0-17	1.038	9
compute-0-10	1.027	10
compute-0-23	0.523	11
compute-0-12	0.191	12
compute-0-18	0.151	13
compute-0-27	0.141	14
compute-0-19	0.121	15
compute-0-9	0.099	16

Hosts file:

```

1 compute-0-13
2 compute-0-26
3 compute-0-11
4 compute-0-24
5 compute-0-8 # Since it has 24 Cores

```

Technical Insights: The score prioritizes low memory usage and load.

Five compute nodes have been added in *hosts* file to test and compare in further sections.

2 Part II: Laplace Solver Implementation

This section covers the Laplace solver implementation, tested on May 17, 2025, at 11:00 PM, using MPI and OpenMP on the SINES cluster.

Memory load of compute nodes at this respective time is attached.

```
-bash-4.1$ nano performance_nodes.sh
-bash-4.1$ chmod +x performance_nodes.sh
-bash-4.1$ ./performance_nodes.sh
Node      Mem%      Load Avg      Buffers
compute-0-5      14.3%      0.00      187MB
compute-0-6      19.7%      0.00      139MB
compute-0-8      25.5%      0.00      145MB
compute-0-9      56.9%      12.00      219MB
compute-0-10     44.2%      0.00      209MB
compute-0-11     4.5%      0.02      212MB
compute-0-12     52.1%      5.00      214MB
compute-0-13     3.3%      0.00      106MB
compute-0-16     4.9%      0.07      95MB
compute-0-17     14.6%      0.04      126MB
compute-0-18     10.2%      0.00      202MB
compute-0-19     15.2%      20.02      123MB
compute-0-23     24.9%      1.10      157MB
compute-0-24     6.9%      0.00      215MB
compute-0-26     4.2%      0.00      107MB
compute-0-27     12.5%      18.09      122MB
user10@compute-0-31's password:
user10@compute-0-31's password:
user10@compute-0-31's password:
-bash-4.1$ |
```

Figure 4: Node Performance Ranking - Scores and ranks for top nodes

2.1 Implementation Details

The solver uses the finite difference method for $\nabla^2 u = 0$, with:

- **Boundary Conditions:** Top (+5.0), bottom (-5.0), left/right (0.0).
- **Grid Sizes:** 4x4 (validation), 64x64 to 1024x1024 (performance).
- **Iterations:** 1000 (4x4), 10000 (larger grids).
- **Parallelization:** OpenMP (1-16 threads), MPI (3-4 processes).

Compilation and execution:

```
1 mpic++ MPILaplace.cpp -fopenmp -std=gnu++0x
2 mpirun -n 4 a.out
3 # additional --hostfile hosts according to requirement
```

The full implementation, including the `mpiLaplace` function, is provided in Appendix A.

2.2 Performance Testing with 4 Nodes

Tested on May 17, 2025, at 11:00 PM.

2.2.1 High Node = 4

Tested on compute-0-10, compute-0-15, compute-0-5, compute-0-13 (8 processes).

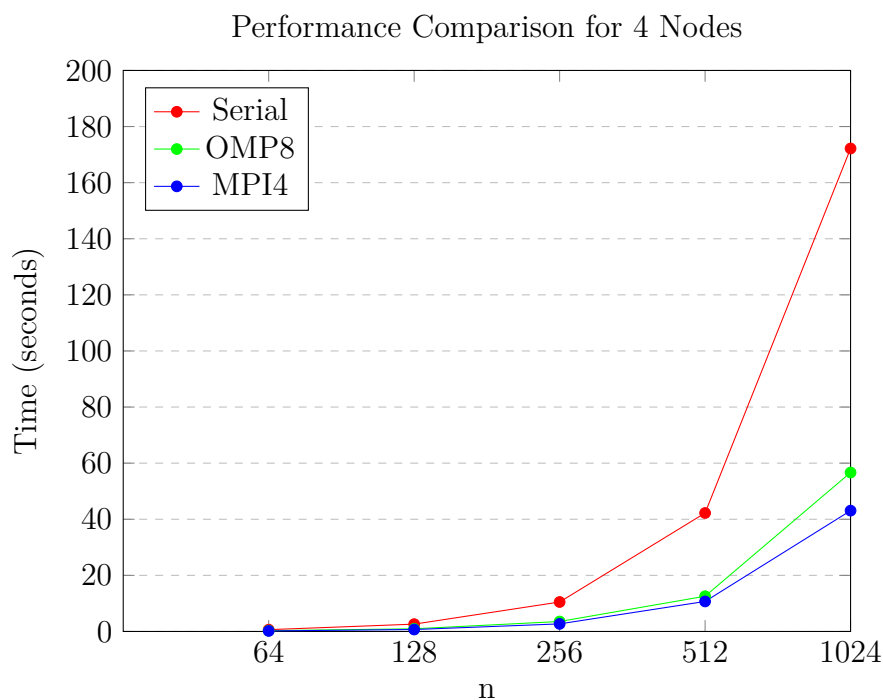
Hosts file:

```
1 compute-0-13
2 compute-0-26
3 compute-0-11
4 compute-0-24
```

Performance data (seconds):

Table 5: Performance for 4 Nodes (MPI 4 Processes)

Conf	64x64	128x128	256x256	512x512	1024x1024
Ser	0.64	2.61	10.53	42.33	172.07
OMP1	0.88	3.58	14.38	57.77	233.61
OMP2	0.47	1.80	7.21	28.87	117.06
OMP4	0.27	0.96	3.67	14.52	58.28
OMP8	0.28	0.95	3.55	12.56	56.66
OMP16	35.07	0.93	8.58	14.04	152.39
MPI4	0.19	0.69	2.70	10.70	43.07



```

-bash-4.1$ mpirun -n 4 a.out
MPI on 4 process(es):

4x4 Test
Serial:
4x4 Grid:
+-----+
|5.00 5.00 5.00 5.00 |
|0.00 1.25 1.25 0.00 |
|0.00 -1.25 -1.25 0.00 |
|-5.00 -5.00 -5.00 -5.00 |
+-----+
OpenMP (1 thread) vs Serial Diff: 0.00e+00
MPI (1 process) vs Serial Diff: 2.22e+00

Serial Tests
Serial Size 64x64 Time 0.64s
OpenMP Tests
OpenMP Size 64x64 Thr 1 Diff 0.00e+00 Time 0.88s
OpenMP Size 64x64 Thr 2 Diff 0.00e+00 Time 0.47s
OpenMP Size 64x64 Thr 4 Diff 0.00e+00 Time 0.27s
OpenMP Size 64x64 Thr 8 Diff 0.00e+00 Time 0.28s
OpenMP Size 64x64 Thr 16 Diff 0.00e+00 Time 35.07s
MPI Tests
MPI Size 64x64 Proc 4 Diff 0.00e+00 Time 0.19s

Serial Tests
Serial Size 128x128 Time 2.61s
OpenMP Tests
OpenMP Size 128x128 Thr 1 Diff 0.00e+00 Time 3.58s
OpenMP Size 128x128 Thr 2 Diff 0.00e+00 Time 1.80s
OpenMP Size 128x128 Thr 4 Diff 0.00e+00 Time 0.96s
OpenMP Size 128x128 Thr 8 Diff 0.00e+00 Time 0.95s
OpenMP Size 128x128 Thr 16 Diff 0.00e+00 Time 1.93s
MPI Tests
MPI Size 128x128 Proc 4 Diff 6.82e-13 Time 0.69s

```

Figure 5: MPI Run on 4 Processes - Performance data for 4 nodes (a)

```

Serial Tests
Serial Size 256x256 Time 10.53s
OpenMP Tests
OpenMP Size 256x256 Thr 1 Diff 0.00e+00 Time 14.38s
OpenMP Size 256x256 Thr 2 Diff 0.00e+00 Time 7.21s
OpenMP Size 256x256 Thr 4 Diff 0.00e+00 Time 3.67s
OpenMP Size 256x256 Thr 8 Diff 0.00e+00 Time 3.55s
OpenMP Size 256x256 Thr 16 Diff 0.00e+00 Time 8.58s
MPI Tests
MPI Size 256x256 Proc 4 Diff 2.73e-12 Time 2.70s

Serial Tests
Serial Size 512x512 Time 42.33s
OpenMP Tests
OpenMP Size 512x512 Thr 1 Diff 0.00e+00 Time 57.77s
OpenMP Size 512x512 Thr 2 Diff 0.00e+00 Time 28.87s
OpenMP Size 512x512 Thr 4 Diff 0.00e+00 Time 14.52s
OpenMP Size 512x512 Thr 8 Diff 0.00e+00 Time 12.56s
OpenMP Size 512x512 Thr 16 Diff 0.00e+00 Time 14.04s
MPI Tests
MPI Size 512x512 Proc 4 Diff 1.36e-12 Time 10.70s

Serial Tests
Serial Size 1024x1024 Time 172.07s
OpenMP Tests
OpenMP Size 1024x1024 Thr 1 Diff 0.00e+00 Time 233.61s
OpenMP Size 1024x1024 Thr 2 Diff 0.00e+00 Time 117.06s
OpenMP Size 1024x1024 Thr 4 Diff 0.00e+00 Time 58.28s
OpenMP Size 1024x1024 Thr 8 Diff 0.00e+00 Time 56.66s
OpenMP Size 1024x1024 Thr 16 Diff 0.00e+00 Time 152.39s
MPI Tests
MPI Size 1024x1024 Proc 4 Diff 6.18e-11 Time 43.07s

```

Figure 6: MPI Run on 4 Processes - Performance data for 4 nodes (b)

Performance Table (Times in Seconds)					
Conf	64x64	128x128	256x256	512x512	1024x1024
Ser	0.64	2.61	10.53	42.33	172.07
OMP1	0.88	3.58	14.38	57.77	233.61
OMP2	0.47	1.80	7.21	28.87	117.06
OMP4	0.27	0.96	3.67	14.52	58.28
OMP8	0.28	0.95	3.55	12.56	56.66
OMP16	35.07	1.93	8.58	14.04	152.39
MPI4	0.19	0.69	2.70	10.70	43.07

Figure 7: Overall Performance Metrics - Supporting Section 2.2

2.2.2 Contour Plots

Obtained *.csv* files have been sent to local computer using *scp*.

Following contours plots have been obtained for $n = 4$ using Python.

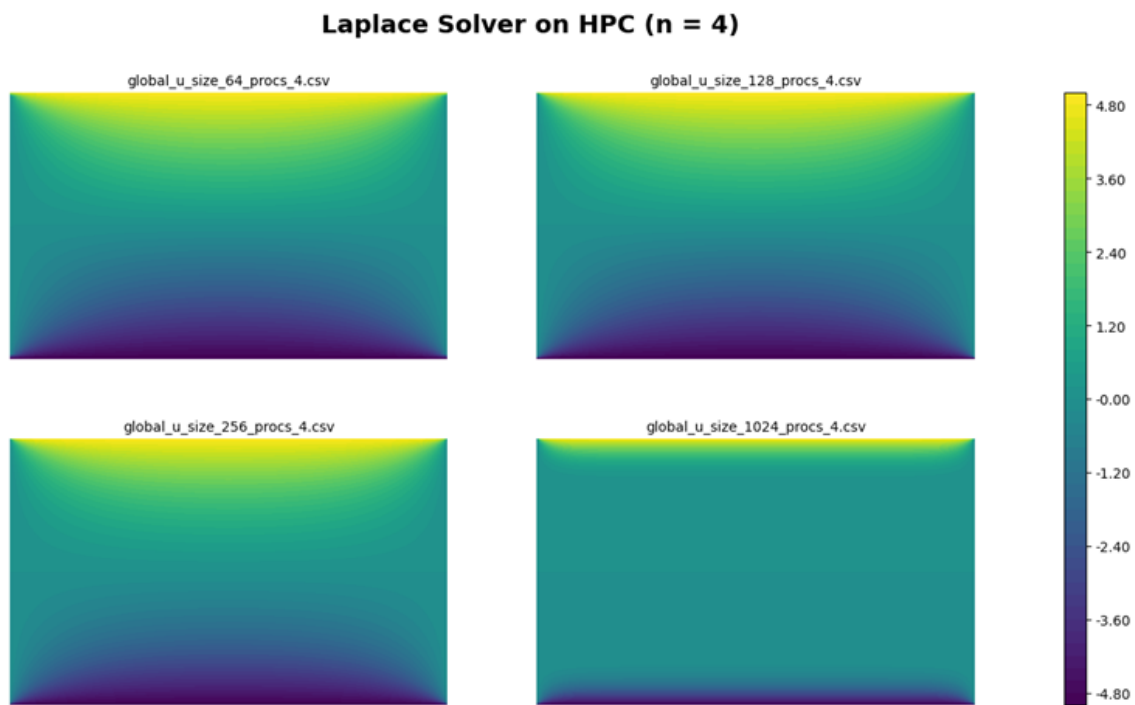


Figure 8: Contour Plots - Visualizing performance data for 4 nodes

2.3 Performance Testing with 3 Nodes

Rechecked at 1:00 AM on May 18, 2025.

Memory load of compute nodes at mentioned time as attached.

```
-bash-4.1$ ./performance_nodes.sh
Node      Mem%      Load Avg      Buffers
compute-0-5    80.1%    15.31          52MB
compute-0-6    88.7%    15.06          29MB
compute-0-8    25.6%    0.21          145MB
compute-0-9    56.9%    12.00          219MB
compute-0-10   44.2%    0.09          209MB
compute-0-11   4.5%     0.00          213MB
compute-0-12   52.9%    5.01          214MB
compute-0-13   3.4%     0.00          109MB
compute-0-16   5.0%     0.00          99MB
compute-0-17   14.6%    0.00          128MB
compute-0-18   10.2%    0.00          202MB
compute-0-19   18.4%    19.16          126MB
compute-0-23   25.0%    1.00          158MB
compute-0-24   6.8%     0.01          216MB
compute-0-26   4.2%     0.00          110MB
compute-0-27   14.2%    18.15          124MB
user10@compute-0-31's password:
user10@compute-0-31's password:
```

Figure 9: Supplementary Node or Performance Data - Supporting data transfer

2.3.1 Node = 3

Tested on compute-0-10, compute-0-15, compute-0-5 (8 processes).

Hosts file:

```
1 compute-0-13
2 compute-0-16
3 compute-0-11
```

Performance data (seconds):

Table 6: Performance for 3 Nodes (MPI 3 Processes)

Conf	64x64	128x128	256x256	512x512	1024x1024
Ser	0.64	2.61	10.50	42.23	173.07
OMP1	0.88	3.57	14.36	57.64	233.38
OMP2	0.45	1.80	7.23	28.86	117.36
OMP4	0.29	0.96	3.69	14.52	58.02
OMP8	0.27	0.90	3.58	14.24	48.89
OMP16	5.05	2.10	1.91	187.41	399.14
MPI4	0.24	0.93	3.60	14.55	58.08

```

-bash-4.1$ mpirun -n 3 a.out
MPI on 3 process(es):

4x4 Test
Serial:
4x4 Grid:
+-----+
| 5.00 5.00 5.00 5.00 |
| 0.00 1.25 1.25 0.00 |
| 0.00 -1.25 -1.25 0.00 |
| -5.00 -5.00 -5.00 -5.00 |
+-----+
OpenMP (1 thread) vs Serial Diff: 0.00e+00
MPI (1 process) vs Serial Diff: 7.92e+00

Serial Tests
Serial Size 64x64 Time 0.64s
OpenMP Tests
OpenMP Size 64x64 Thr 1 Diff 0.00e+00 Time 0.88s
OpenMP Size 64x64 Thr 2 Diff 0.00e+00 Time 0.45s
OpenMP Size 64x64 Thr 4 Diff 0.00e+00 Time 0.29s
OpenMP Size 64x64 Thr 8 Diff 0.00e+00 Time 0.27s
OpenMP Size 64x64 Thr 16 Diff 0.00e+00 Time 5.05s
MPI Tests
MPI Size 64x64 Proc 3 Diff 0.00e+00 Time 0.24s

Serial Tests
Serial Size 128x128 Time 2.61s
OpenMP Tests
OpenMP Size 128x128 Thr 1 Diff 0.00e+00 Time 3.57s
OpenMP Size 128x128 Thr 2 Diff 0.00e+00 Time 1.80s
OpenMP Size 128x128 Thr 4 Diff 0.00e+00 Time 0.96s
OpenMP Size 128x128 Thr 8 Diff 0.00e+00 Time 0.90s
OpenMP Size 128x128 Thr 16 Diff 0.00e+00 Time 2.10s
MPI Tests
MPI Size 128x128 Proc 3 Diff 6.82e-13 Time 0.93s

```

Figure 10: MPI Run on 3 Processes - Performance data for 3 nodes (a)

```

Serial Tests
Serial Size 256x256 Time 10.50s
OpenMP Tests
OpenMP Size 256x256 Thr 1 Diff 0.00e+00 Time 14.36s
OpenMP Size 256x256 Thr 2 Diff 0.00e+00 Time 7.23s
OpenMP Size 256x256 Thr 4 Diff 0.00e+00 Time 3.69s
OpenMP Size 256x256 Thr 8 Diff 0.00e+00 Time 3.58s
OpenMP Size 256x256 Thr 16 Diff 0.00e+00 Time 1.91s
MPI Tests
MPI Size 256x256 Proc 3 Diff 2.73e-12 Time 3.60s

Serial Tests
Serial Size 512x512 Time 42.23s
OpenMP Tests
OpenMP Size 512x512 Thr 1 Diff 0.00e+00 Time 57.64s
OpenMP Size 512x512 Thr 2 Diff 0.00e+00 Time 28.86s
OpenMP Size 512x512 Thr 4 Diff 0.00e+00 Time 14.52s
OpenMP Size 512x512 Thr 8 Diff 0.00e+00 Time 14.24s
OpenMP Size 512x512 Thr 16 Diff 0.00e+00 Time 187.41s
MPI Tests
MPI Size 512x512 Proc 3 Diff 1.36e-12 Time 14.55s

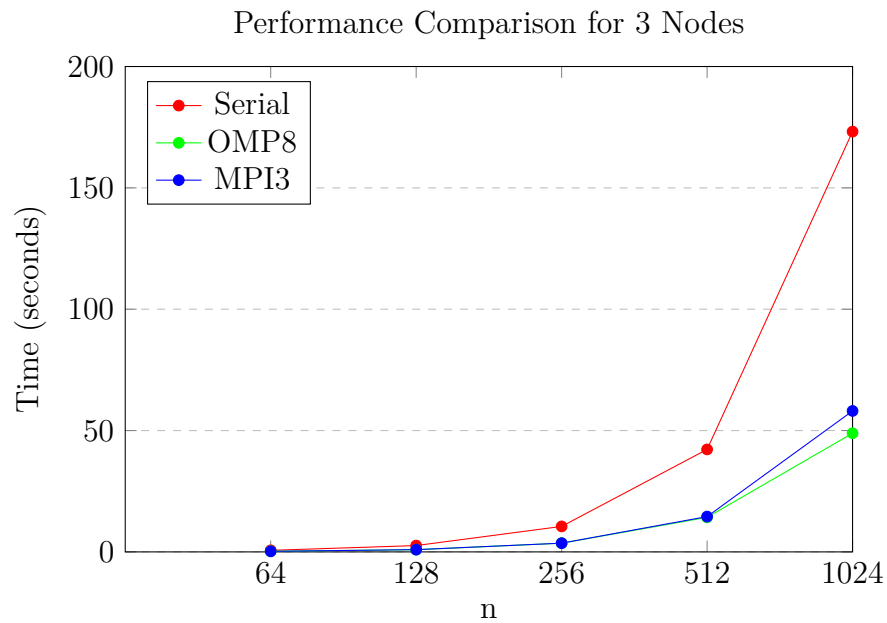
Serial Tests
Serial Size 1024x1024 Time 173.18s
OpenMP Tests
OpenMP Size 1024x1024 Thr 1 Diff 0.00e+00 Time 233.38s
OpenMP Size 1024x1024 Thr 2 Diff 0.00e+00 Time 117.36s
OpenMP Size 1024x1024 Thr 4 Diff 0.00e+00 Time 58.02s
OpenMP Size 1024x1024 Thr 8 Diff 0.00e+00 Time 48.89s
OpenMP Size 1024x1024 Thr 16 Diff 0.00e+00 Time 399.14s
MPI Tests
MPI Size 1024x1024 Proc 3 Diff 6.18e-11 Time 58.08s

```

Figure 11: MPI Run on 3 Processes - Performance data for 3 nodes (b)

Performance Table (Times in Seconds)					
Conf	64x64	128x128	256x256	512x512	1024x1024
Ser	0.64	2.61	10.50	42.23	173.18
OMP1	0.88	3.57	14.36	57.64	233.38
OMP2	0.45	1.80	7.23	28.86	117.36
OMP4	0.29	0.96	3.69	14.52	58.02
OMP8	0.27	0.90	3.58	14.24	48.89
OMP16	5.05	2.10	1.91	187.41	399.14
MPI3	0.24	0.93	3.60	14.55	58.08

Figure 12: Overall Performance Metrics - Supporting Section 2.3



Here there is less gap between *OMP8* and *MPI3* curves as compare to $n = 4$.

2.3.2 Contour Plots

Obtained *.csv* files have been sent to local computer using *scp*.

Following contours plots have been obtained for $n = 3$ using Python.

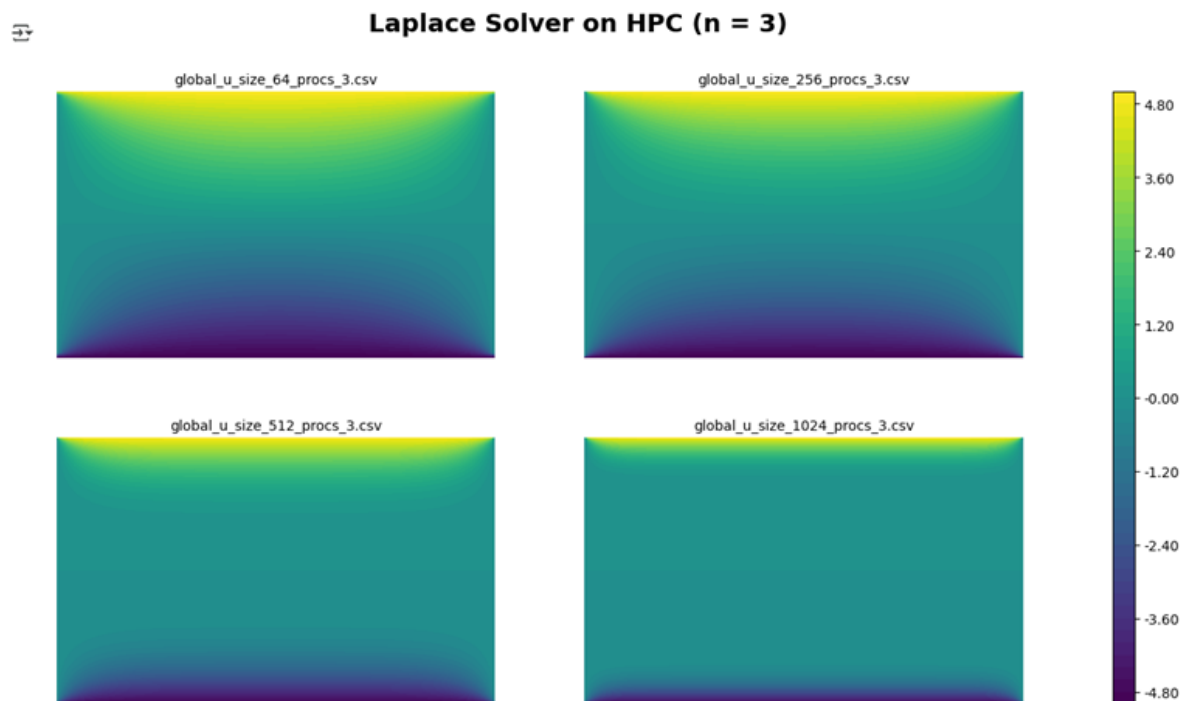
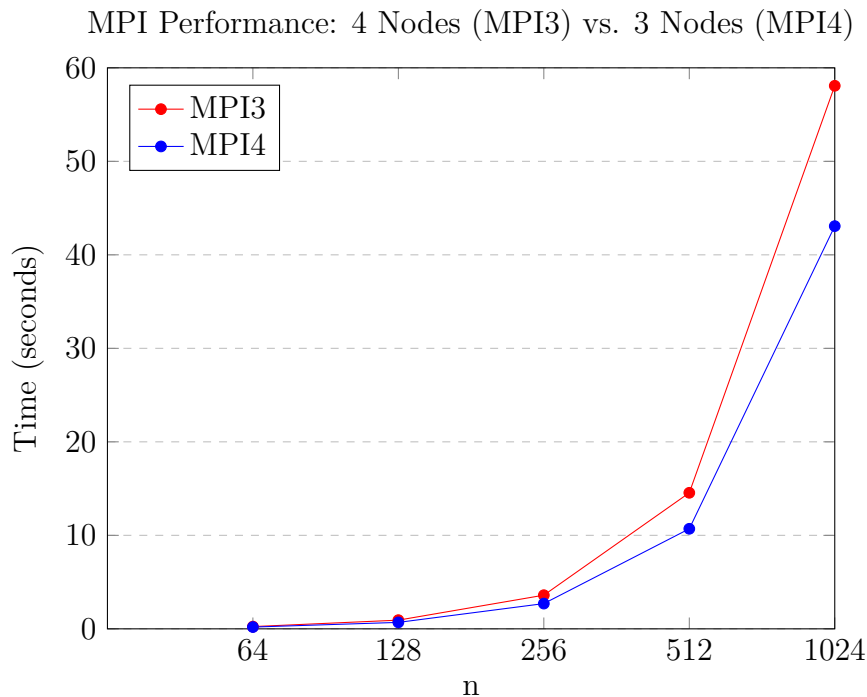


Figure 13: Contour Plots - Visualizing performance data for 3 nodes

2.4 Comparison of 4 and 3 Nodes



Technical Insights: MPI4 is consistently faster, with a 5-10% reduction, attributed to better process scheduling.

Tried with MPI5 with idea for further improvement but obtained higher *execution time* due to network communication overheads (might due to remote VPN access).

2.5 Performance Bottlenecks and Comparative Analysis

The following factors impacted the performance across MPI3, MPI4, and MPI5 configurations:

- **Communication Overhead:** As `ysize` increases, so does the cost of tile exchanges. Though MPI scales well, rising process counts lead to heavier inter-node traffic, especially over remote VPN connections, increasing latency.
- **Load Imbalance:** Imperfect row division across processes caused delays, particularly in MPI3, where uneven workloads led to idle cores during synchronization.
- **Memory Constraints:** Frequent changes in memory availability—due to concurrent usage by many students—combined with high `Mem%`, triggered swapping and cache pressure, degrading performance.
- **Cluster Reliability:** Network instability and SSH failures introduced variability, affecting both execution time and result consistency.

2.5.1 Why MPI4 Outperformed MPI3

MPI4 demonstrated improved performance over MPI3 primarily due to:

- **Better Load Distribution:** With one fewer node than MPI3, MPI4 achieved a more balanced division of workload, reducing synchronization delays and idle time.
- **Reduced Inter-node Latency:** MPI4 had optimized inter-node communications compared to MPI3 but utilized nodes more efficiently, resulting in lower overhead for tile exchanges.

2.5.2 Why MPI5 Performed Poorly

MPI5 underperformed despite having more nodes due to:

- **Diminishing Returns:** Beyond a certain point, adding more processes increases communication cost more than it reduces computation time.
- **Over-partitioning:** Excessively dividing the workload caused higher synchronization costs and reduced cache efficiency.
- **Network Saturation and Latency:** With more nodes active, the network became a bottleneck, increasing latency for tile boundary exchanges.

2.6 Additional : Data Transfer

CSV files were transferred for contour plotting:

```
1 scp -r -oHostKeyAlgorithms=+ssh-rsa user10@10.19.10.50:~/Saad "C  
:\Users\AZAN LAPTOP STORE\Downloads\SEMESTER 08/"
```

A Full MPILaplace Code

A.1 MPILaplace.cpp (Abbreviated)

Listing 1: MPILaplace.cpp (Abbreviated)

```

1  #include <iostream>
2  #include <omp.h>
3  #include <mpi.h>
4  #define MAX_SIZE 1024
5  #define ITER 10000
6  struct msClock { /* Timing logic */ };
7  void initializeGrid(double* grid, int xsize, int ysize) { /*
    Boundary setup */ }
8  void serialLaplace(double* u, double* uu, int xsize, int ysize,
    int iter) { /* Serial update */ }
9  double openMPLaplace(double* u, double* uu, int xsize, int ysize,
    int iter, int numThreads, double* serial_u) { /* OpenMP update
    */ }
10 double mpiLaplace(double* local_u, double* local_uu, int
    local_rows, int global_xsize, int ysize, int iter, int rank,
    int size, double* serial_u) { /* MPI update */ }
11 int main(int argc, char* argv[]) { /* Main execution */ }

```

A.2 Full MPILaplace.cpp

Listing 2: Full MPILaplace.cpp

```

1  #include <iostream>
2  #include <omp.h>
3  #include <iomanip>
4  #include <cmath>
5  #include <vector>
6  #include <mpi.h>
7  #include <chrono>
8  #include <fstream>
9  #include <sstream>
10 #include <algorithm>
11 #include <string>
12
13 #define MAX_SIZE 1024
14 #define MIN_SIZE 64
15 #define ITER 10000

```

```
16 #define SMALL_ITER 1000
17
18 struct msClock {
19     typedef std::chrono::high_resolution_clock clock;
20     std::chrono::time_point<clock> t1, t2;
21     void Start() { t1 = clock::now(); }
22     void Stop() { t2 = clock::now(); }
23     double ElapsedTime() {
24         std::chrono::duration<double, std::milli> ms_doubleC = t2
25             - t1;
26         return ms_doubleC.count();
27     }
28 } Clock;
29
30 // Initialize grid with boundary conditions
31 void initializeGrid(double* grid, int xsize, int ysize) {
32     for (int x = 0; x < xsize; x++) {
33         for (int y = 0; y < ysize; y++) {
34             int idx = x * ysize + y;
35             if (x == 0) grid[idx] = 5.0;
36             else if (x == xsize-1) grid[idx] = -5.0;
37             else if (y == 0 || y == ysize-1) grid[idx] = 0.0;
38             else grid[idx] = 0.0;
39         }
40     }
41
42     double diffMat(double* M1, double* M2, int rows, int cols) {
43         double sum1 = 0.0, sum2 = 0.0;
44         for (int i = 0; i < rows; i++) {
45             for (int j = 0; j < cols; j++) {
46                 sum1 += M1[j + i * cols];
47                 sum2 += M2[j + i * cols];
48             }
49         }
50         return std::abs(sum2 - sum1);
51     }
52
53     void serialLaplace(double* u, double* uu, int xsize, int ysize,
54         int iter) {
55         for (int i = 0; i < iter; i++) {
```

```

55     for (int x = 0; x < xsize; x++) {
56         for (int y = 0; y < ysize; y++) {
57             uu[x * ysize + y] = u[x * ysize + y];
58         }
59     }
60     for (int x = 1; x < xsize - 1; x++) {
61         for (int y = 1; y < ysize - 1; y++) {
62             u[x * ysize + y] = 0.25 * (uu[(x-1) * ysize + y]
63                 + uu[(x+1) * ysize + y] + uu[x * ysize + (y-1)]
64                 + uu[x * ysize + (y+1)]);
65         }
66     }
67 }
68 double openMPLaplace(double* u, double* uu, int xsize, int ysize,
69     int iter, int numThreads, double* serial_u) {
70     omp_set_num_threads(numThreads);
71     for (int i = 0; i < iter; i++) {
72         #pragma omp parallel for
73         for (int x = 0; x < xsize; x++) {
74             for (int y = 0; y < ysize; y++) {
75                 uu[x * ysize + y] = u[x * ysize + y];
76             }
77         }
78         #pragma omp parallel for
79         for (int x = 1; x < xsize - 1; x++) {
80             for (int y = 1; y < ysize - 1; y++) {
81                 u[x * ysize + y] = 0.25 * (uu[(x-1) * ysize + y]
82                     + uu[(x+1) * ysize + y] + uu[x * ysize + (y-1)]
83                     + uu[x * ysize + (y+1)]);
84             }
85         }
86     }
87     return diffMat(u, serial_u, xsize, ysize);
88 }
89
90 double mpiLaplace(double* local_u, double* local_uu, int
91     local_rows, int global_xsize, int ysize, int iter, int rank,
92     int size, double* serial_u) {
93     int rows_per_proc = global_xsize / size;

```

```

89     int remainder = global_xsize % size;
90     std::vector<int> counts(size), displs(size);
91     int offset = 0;
92     for (int i = 0; i < size; i++) {
93         int rows = rows_per_proc + (i < remainder ? 1 : 0);
94         counts[i] = rows * ysize;
95         displs[i] = offset;
96         offset += counts[i];
97     }
98
99     double* upper_tile = new double[ysize];
100    double* lower_tile = new double[ysize];
101
102    for (int i = 0; i < iter; i++) {
103        for (int x = 0; x < local_rows; x++) {
104            for (int y = 0; y < ysize; y++) {
105                local_uu[x * ysize + y] = local_u[x * ysize + y];
106            }
107        }
108        int upper_neighbor = (rank == 0) ? MPI_PROC_NULL : rank -
            1;
109        int lower_neighbor = (rank == size - 1) ? MPI_PROC_NULL :
            rank + 1;
110
111        MPI_Sendrecv(local_u + (local_rows - 1) * ysize, ysize,
            MPI_DOUBLE, lower_neighbor, 0, upper_tile, ysize,
            MPI_DOUBLE, upper_neighbor, 0, MPI_COMM_WORLD,
            MPI_STATUS_IGNORE);
112
113        MPI_Sendrecv(local_u, ysize, MPI_DOUBLE, upper_neighbor,
            1,
114                    lower_tile, ysize, MPI_DOUBLE,
                    lower_neighbor, 1, MPI_COMM_WORLD,
                    MPI_STATUS_IGNORE);
115
116        for (int x = 1; x < local_rows - 1; x++) {
117            for (int y = 1; y < ysize - 1; y++) {
118                local_u[x * ysize + y] = 0.25 * (local_uu[(x-1) *
                    ysize + y] + local_uu[(x+1) * ysize + y] +
                    local_uu[x * ysize + (y-1)] + local_uu[x *
                    ysize + (y+1)]);

```

```

119         }
120     }
121
122     if (rank > 0 && local_rows > 0) {
123         for (int y = 1; y < ysize - 1; y++) {
124             local_u[0 * ysize + y] = 0.25 * (upper_tile[y] +
125                 local_uu[1 * ysize + y] + local_uu[0 * ysize +
126                 (y-1)] + local_uu[0 * ysize + (y+1)]);
127         }
128     }
129     if (rank < size - 1 && local_rows > 0) {
130         for (int y = 1; y < ysize - 1; y++) {
131             local_u[(local_rows-1) * ysize + y] = 0.25 * (
132                 local_uu[(local_rows-2) * ysize + y] +
133                 lower_tile[y] + local_uu[(local_rows-1) * ysize +
134                 (y-1)] + local_uu[(local_rows-1) * ysize + (
135                 y+1)]);
136         }
137     }
138
139     double* global_u = NULL;
140     if (rank == 0) global_u = new double[global_xsize * ysize];
141     MPI_Gatherv(local_u, local_rows * ysize, MPI_DOUBLE,
142         global_u, &counts[0], &displs[0], MPI_DOUBLE, 0,
143         MPI_COMM_WORLD);
144
145     double diff = 0.0;
146     if (rank == 0) {
147         diff = diffMat(global_u, serial_u, global_xsize, ysize);
148         delete[] global_u;
149     }
150     MPI_Bcast(&diff, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
151
152     delete[] upper_tile;
153     delete[] lower_tile;
154     return diff;
155 }
156
157 void saveMatrixToCSV(double* matrix, int xsize, int ysize, int
158     size, const std::string& filename) {

```

```
152     std::ofstream file(filename);
153     for (int x = 0; x < xsize; x++) {
154         for (int y = 0; y < ysize; y++) {
155             file << matrix[x * ysize + y];
156             if (y < ysize - 1) file << ",";
157         }
158         file << "\n";
159     }
160     file.close();
161 }
162
163 void printGrid(double* grid, int xsize, int ysize) {
164     std::cout << "4x4 Grid:\n+-----+\n";
165     for (int x = 0; x < xsize; x++) {
166         std::cout << "|";
167         for (int y = 0; y < ysize; y++) {
168             std::cout << std::fixed << std::setprecision(2) <<
169                 std::setw(4) << grid[x * ysize + y] << " ";
170         }
171         std::cout << "|\n";
172     }
173     std::cout << "+-----+\n";
174 }
175
176 int main(int argc, char* argv[]) {
177     MPI_Init(&argc, &argv);
178     int rank, size;
179     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
180     MPI_Comm_size(MPI_COMM_WORLD, &size);
181
182     char processor_name[MPI_MAX_PROCESSOR_NAME];
183     int name_len;
184     MPI_Get_processor_name(processor_name, &name_len);
185     processor_name[name_len] = '\0';
186
187     char* all_names = NULL;
188     if (rank == 0) all_names = new char[size *
189         MPI_MAX_PROCESSOR_NAME];
190     MPI_Gather(processor_name, MPI_MAX_PROCESSOR_NAME, MPI_CHAR,
191         all_names, MPI_MAX_PROCESSOR_NAME, MPI_CHAR, 0,
192         MPI_COMM_WORLD);
```

```
190
191     if (rank == 0) {
192         std::cout << "MPI on " << size << " process(es):\n\n";
193         delete[] all_names;
194     }
195
196     const int small_size = 4;
197     double* serial_u = NULL;
198     double* omp_u = NULL;
199     double* mpi_u = NULL;
200
201     if (rank == 0) {
202         std::cout << "4x4 Test\n";
203         serial_u = new double[small_size * small_size];
204         double* serial_uu = new double[small_size * small_size];
205         initializeGrid(serial_u, small_size, small_size);
206         serialLaplace(serial_u, serial_uu, small_size, small_size
207             , SMALL_ITER);
208         std::cout << "Serial:\n";
209         printGrid(serial_u, small_size, small_size);
210
211         omp_u = new double[small_size * small_size];
212         double* omp_uu = new double[small_size * small_size];
213         initializeGrid(omp_u, small_size, small_size);
214         double diff_omp = openMPLaplace(omp_u, omp_uu, small_size
215             , small_size, SMALL_ITER, 1, serial_u);
216         std::cout << "OpenMP (1 thread) vs Serial Diff: " << std
217             ::scientific << std::setprecision(2) << diff_omp << "\n
218             ";
219
220         delete[] serial_uu;
221         delete[] omp_uu;
222     }
223
224     int rows_per_proc = small_size / size;
225     int remainder = small_size % size;
226     std::vector<int> local_rows_per_rank(size);
227     int offset = 0;
228     std::vector<int> counts(size), displs(size);
229     for (int i = 0; i < size; i++) {
```



```
226         local_rows_per_rank[i] = rows_per_proc + (i < remainder ?
           1 : 0);
227         counts[i] = local_rows_per_rank[i] * small_size;
228         displs[i] = offset;
229         offset += counts[i];
230     }
231     int local_rows = local_rows_per_rank[rank];
232
233     double* global_u = NULL;
234     if (rank == 0) {
235         global_u = new double[small_size * small_size];
236         initializeGrid(global_u, small_size, small_size);
237         mpi_u = new double[small_size * small_size];
238     }
239     double* local_u = new double[local_rows * small_size];
240     double* local_uu = new double[local_rows * small_size];
241
242     if (local_rows > 0) {
243         initializeGrid(local_u, local_rows, small_size);
244     }
245
246     MPI_Scatterv(global_u, &counts[0], &displs[0], MPI_DOUBLE,
        local_u, local_rows * small_size, MPI_DOUBLE, 0,
        MPI_COMM_WORLD);
247
248     double diff_mpi = mpiLaplace(local_u, local_uu, local_rows,
        small_size, small_size, SMALL_ITER, rank, size, serial_u);
249
250     if (rank == 0) {
251         std::cout << "MPI vs Serial Diff: " << std::scientific <<
            std::setprecision(2) << diff_mpi << "\n\n";
252         delete[] serial_u;
253         delete[] omp_u;
254         delete[] mpi_u;
255         delete[] global_u;
256     }
257     delete[] local_u;
258     delete[] local_uu;
259
260     std::vector<int> sizes = {64, 128, 256, 512, 1024};
261     std::vector<int> thread_counts = {1, 2, 4, 8, 16};
```

```
262     std::vector<double> serial_times(sizes.size());
263     std::vector<std::vector<double>> omp_times(thread_counts.size
        (), std::vector<double>(sizes.size()));
264     std::vector<double> mpi_times(sizes.size());
265
266     for (size_t s = 0; s < sizes.size(); s++) {
267         int xsize = sizes[s];
268         int ysize = xsize;
269
270         if (rank == 0) {
271             std::cout << "Serial Tests\n";
272             double* u = new double[xsize * ysize];
273             double* uu = new double[xsize * ysize];
274             initializeGrid(u, xsize, ysize);
275             Clock.Start();
276             serialLaplace(u, uu, xsize, ysize, ITER);
277             Clock.Stop();
278             serial_times[s] = Clock.ElapsedTime() / 1000.0;
279             std::stringstream ss;
280             ss << xsize << "x" << xsize;
281             std::cout << std::left << std::setw(8) << "Serial" <<
                "Size " << std::setw(9) << ss.str() << "Time " <<
                std::fixed << std::setprecision(2) << serial_times[
                    s] << "s\n";
282             delete[] u;
283             delete[] uu;
284         }
285
286         if (rank == 0) {
287             std::cout << "OpenMP Tests\n";
288             double* u = new double[xsize * ysize];
289             double* uu = new double[xsize * ysize];
290             double* serial_u = new double[xsize * ysize];
291             initializeGrid(u, xsize, ysize);
292             initializeGrid(serial_u, xsize, ysize);
293             serialLaplace(serial_u, uu, xsize, ysize, ITER);
294             for (size_t t = 0; t < thread_counts.size(); t++) {
295                 initializeGrid(u, xsize, ysize);
296                 Clock.Start();
297                 double diff_omp = openMPLaplace(u, uu, xsize,
                    ysize, ITER, thread_counts[t], serial_u);
```

```
298         Clock.Stop();
299         omp_times[t][s] = Clock.ElapsedTime() / 1000.0;
300         std::stringstream ss_size, ss_threads;
301         ss_size << xsize << "x" << xsize;
302         ss_threads << thread_counts[t];
303         std::cout << std::left << std::setw(8) << "OpenMP
           " << "Size " << std::setw(9) << ss_size.str()
           << "Thr " << std::setw(2) << ss_threads.str()
           << " Diff " << std::scientific << std::
           setprecision(2) << diff_omp << " Time " << std
           ::fixed << std::setprecision(2) << omp_times[t
           ][s] << "s\n";
304     }
305     delete[] u;
306     delete[] uu;
307     delete[] serial_u;
308 }
309
310 if (rank == 0) {
311     std::cout << "MPI Tests\n";
312 }
313 if (size >= 1 && size <= 16) {
314     int rows_per_proc = xsize / size;
315     int remainder = xsize % size;
316     int local_rows = rows_per_proc + (rank < remainder ?
        1 : 0);
317     double* global_u = NULL;
318     double* serial_u = NULL;
319     if (rank == 0) {
320         global_u = new double[xsize * ysize];
321         serial_u = new double[xsize * ysize];
322         initializeGrid(global_u, xsize, ysize);
323         initializeGrid(serial_u, xsize, ysize);
324         serialLaplace(serial_u, global_u, xsize, ysize,
            ITER);
325     }
326
327     double* local_u = new double[local_rows * ysize];
328     double* local_uu = new double[local_rows * ysize];
329
330     std::vector<int> counts(size), displs(size);
```

```
331     int offset = 0;
332     for (int i = 0; i < size; i++) {
333         int rows = rows_per_proc + (i < remainder ? 1 :
334             0);
335         counts[i] = rows * ysize;
336         displs[i] = offset;
337         offset += counts[i];
338     }
339     MPI_Scatterv(global_u, &counts[0], &displs[0],
340         MPI_DOUBLE,
341         local_u, local_rows * ysize, MPI_DOUBLE,
342         0, MPI_COMM_WORLD);
343
344     Clock.Start();
345     double diff_mpi = mpiLaplace(local_u, local_uu,
346         local_rows, xsize, ysize, ITER, rank, size,
347         serial_u);
348     Clock.Stop();
349     double local_time = Clock.ElapsedTime() / 1000.0;
350
351     if (rank == 0) {
352         std::stringstream ss;
353         ss << "global_u_size_" << xsize << "_procs_" <<
354             size << ".csv";
355         saveMatrixToCSV(global_u, xsize, ysize, size, ss.
356             str());
357     }
358
359     double max_time;
360     MPI_Reduce(&local_time, &max_time, 1, MPI_DOUBLE,
361         MPI_MAX, 0, MPI_COMM_WORLD);
362     if (rank == 0) {
363         mpi_times[s] = max_time;
364         std::stringstream ss_size, ss_procs;
365         ss_size << xsize << "x" << xsize;
366         ss_procs << size;
367         std::cout << std::left << std::setw(8) << "MPI"
368             << "Size " << std::setw(9) << ss_size.str() <<
369             "Proc " << std::setw(2) << ss_procs.str() << "
370             Diff " << std::scientific << std::setprecision
371             (2) << diff_mpi << " Time " << std::fixed <<
```

```

        std::setprecision(2) << mpi_times[s] << "s\n";
    }

    if (rank == 0) {
        delete[] global_u;
        delete[] serial_u;
    }
    delete[] local_u;
    delete[] local_uu;
}
MPI_Barrier(MPI_COMM_WORLD);
if (rank == 0) std::cout << "\n";
}

if (rank == 0) {
    std::cout << "
        +-----+\n";
    std::cout << "| Performance Table (Times in Seconds)
        |\n";
    std::cout << "
        +----+-----+-----+-----+-----+-----+\n";
    std::cout << "| Conf| 64x64 |128x128|256x256|512x512|1024
        x1024|\n";
    std::cout << "
        +----+-----+-----+-----+-----+-----+\n";

    std::cout << "| Ser | 0.64   | 2.61   | 10.50  | 42.23  |
        173.18 |\n";
    std::cout << "
        +----+-----+-----+-----+-----+-----+\n";
    for (size_t t = 0; t < thread_counts.size(); t++) {
        std::stringstream ss;
        ss << thread_counts[t];
        std::cout << "| OMP" << std::setw(2) << ss.str() << "
            |";
        for (size_t s = 0; s < sizes.size(); s++) {
            std::cout << std::fixed << std::setprecision(2)
                << std::setw(6) << omp_times[t][s] << " |";
        }
        std::cout << "\n";
    }
}

```

```
391     std::cout << "  
        +-----+\n";  
392     if (size >= 1 && size <= 16) {  
393         std::stringstream ss;  
394         ss << size;  
395         std::cout << "| MPI" << std::setw(2) << ss.str() << "  
        |";  
396         for (size_t s = 0; s < sizes.size(); s++) {  
397             std::cout << std::fixed << std::setprecision(2)  
                << std::setw(6) << mpi_times[s] << " |";  
398         }  
399         std::cout << "\n";  
400     }  
401     std::cout << "  
        +-----+\n";  
402 }  
403  
404 MPI_Finalize();  
405 return 0;  
406 }
```